

PERFORMANCE OF RDF QUERY PROCESSING ON THE INTEL SCC

Vasil Slavov, Praveen Rao, Dinesh Barenkala, Srivenu Paturi

Department of Computer Science & Electrical Engineering
University of Missouri-Kansas City, USA

Acknowledgements

Single-chip Cloud Computer Research Program, Intel Labs
National Science Foundation (IIS-1115871)

Overview

- ◆ Introduction
- ◆ Background
 - ◆ RDF and SPARQL
 - ◆ Prior work on RDF query processing
- ◆ Our methodology
- ◆ Performance evaluation
- ◆ Conclusion and future work

Introduction (1/2)

◆ The RDF data model

- ◆ Important in domain-specific applications and the WWW
 - ◆ Knowledge representation and reasoning
 - ◆ E.g., healthcare, defense and intelligence, biopharmaceuticals, Linked Data

◆ Very large RDF datasets are available today

- ◆ DBPedia [[WWW '07](#), [ISWC '07](#)]
- ◆ YAGO2 [[WWW '11](#)]
- ◆ Billion Triples Challenge (<http://challenge.semanticweb.org>)
- ◆ Uniprot RDF (<http://www.uniprot.org/downloads>)

Introduction (2/2)

- ◆ Pressing need for high performance RDF processing tools



Can a manycore processor boost the performance of RDF query processing through parallel processing?

Background (1/3)

◆ RDF

- ◆ Triple format: (subject, predicate, object)
- ◆ Represents a directed, labeled graph

Example

```
<Intel_Corporation> rdf:type <wikicategory_Companies_established_in_1968> .  
<Intel_Corporation> rdf:type <wikicategory_Motherboard_companies> .  
<Intel_Corporation> rdf:type <wikicategory_Multinational_companies> .  
<Intel_Corporation> rdf:type <wikicategory_Netbook_manufacturers> .  
<Intel_Corporation> rdf:type <wikicategory_Semiconductor_companies> .  
<Intel_Corporation> y:created <IA-32_Execution_Layer> .  
<Intel_Corporation> y:created <Itanium> .  
<Intel_Corporation> y:created <Light_Peak> .
```

Background (2/3)

◆ SPARQL

- ◆ A popular query language for RDF

Example

@prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .

@prefix y: <http://www.mpii.de/yago/resource/> .

SELECT ?a ?n WHERE

{

?a rdf:type <wikicategory_Motherboard_companies> .

?a y:created ?n .

}

| ?a | ?n |
|---------------------|-------------------------|
| <Intel_Corporation> | <IA-32_Execution_Layer> |
| <Intel_Corporation> | <Itanium> |
| <Intel_Corporation> | <Light_Peak> |

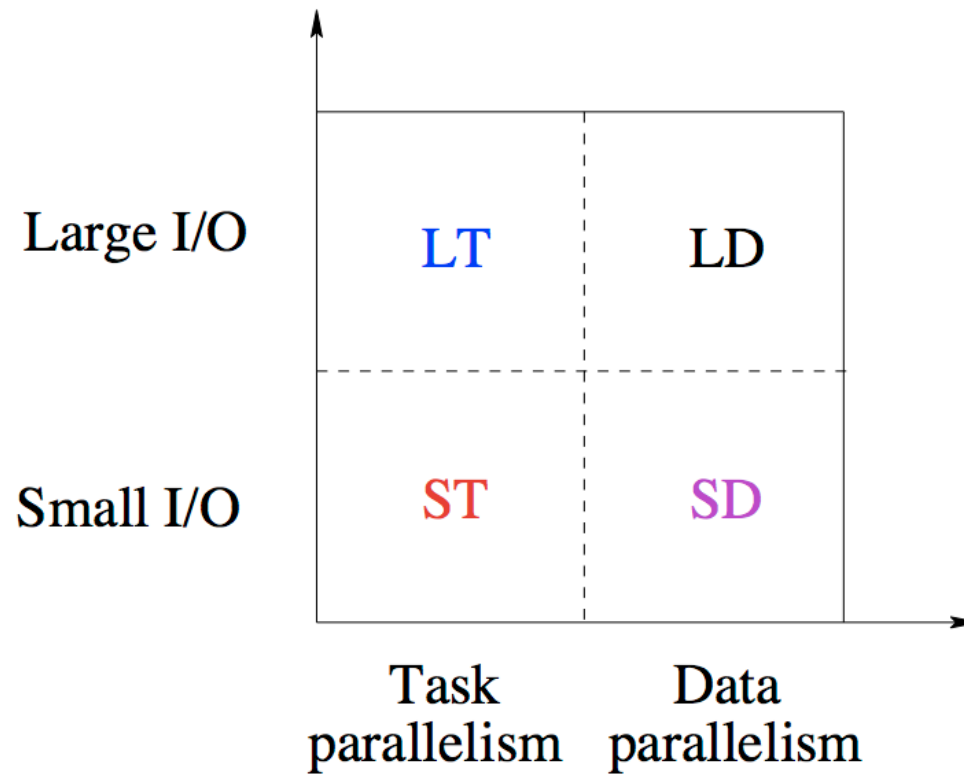
Background (3/3)

- ◆ Prior work on RDF query processing
 - ◆ Using an RDBMS
 - ◆ Sesame [[ISWC '02](#)]
 - ◆ Jena2 [[SWDB '03](#)]
 - ◆ RDF_MATCH [[VLDB '05](#)]
 - ◆ Vertical partitioning [[VLDB '07](#)]
 - ◆ Native RDF databases
 - ◆ Hexastore [[VLDB '08](#)]
 - ◆ RDF-3X [[VLDB '08](#), [VLDB Journal '10](#)]
 - ◆ BitMat [[WWW '10](#)]
 - ◆ Shared-nothing clusters
 - ◆ YARS2 [[ISWC '07](#)], 4store
 - ◆ Tools built using Apache Hadoop and Apache Pig
- ◆ On the Intel SCC (but not RDF)
 - ◆ Parallel AI planning [[MARC '11](#)]
 - ◆ Relational decision support queries [[MARC '11](#)]

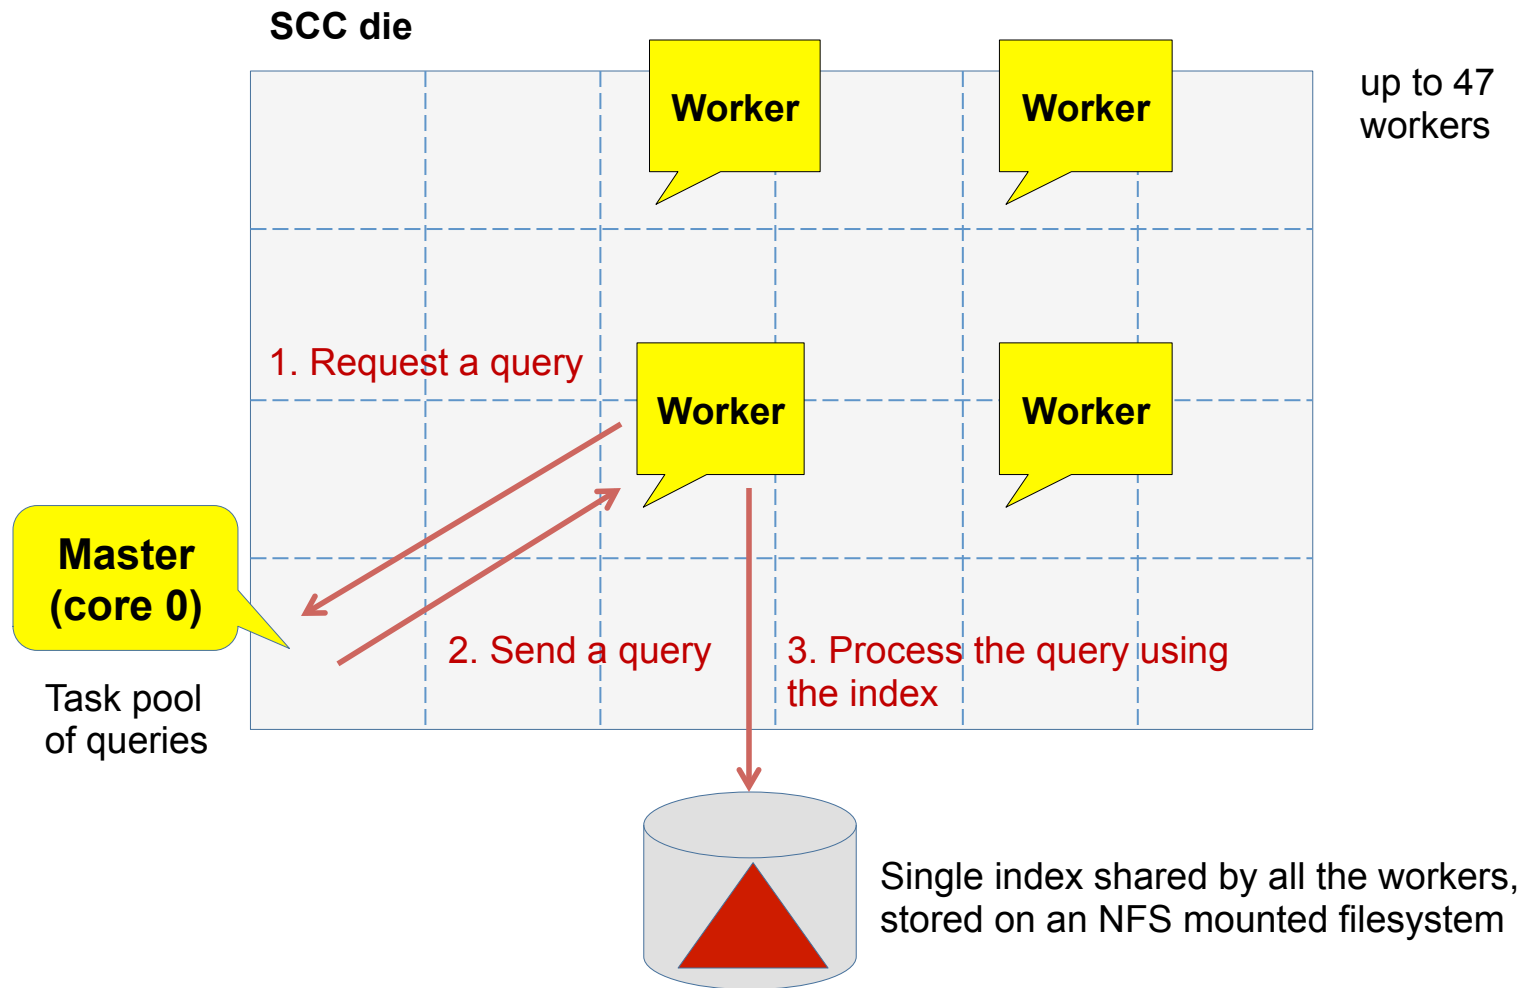
Our Methodology

- ◆ Programming models
 - ◆ Task (inter-query) parallelism
 - ◆ Data (intra-query) parallelism
- ◆ I/O bound queries
 - ◆ Small I/O footprint
 - ◆ Large I/O footprint
- ◆ Message Passing Interface (MPI)
 - ◆ MPI_Send, MPI_Recv
 - ◆ MPI_Barrier
 - ◆ MPI_Bcast, MPI_Scatter, MPI_Gatherv

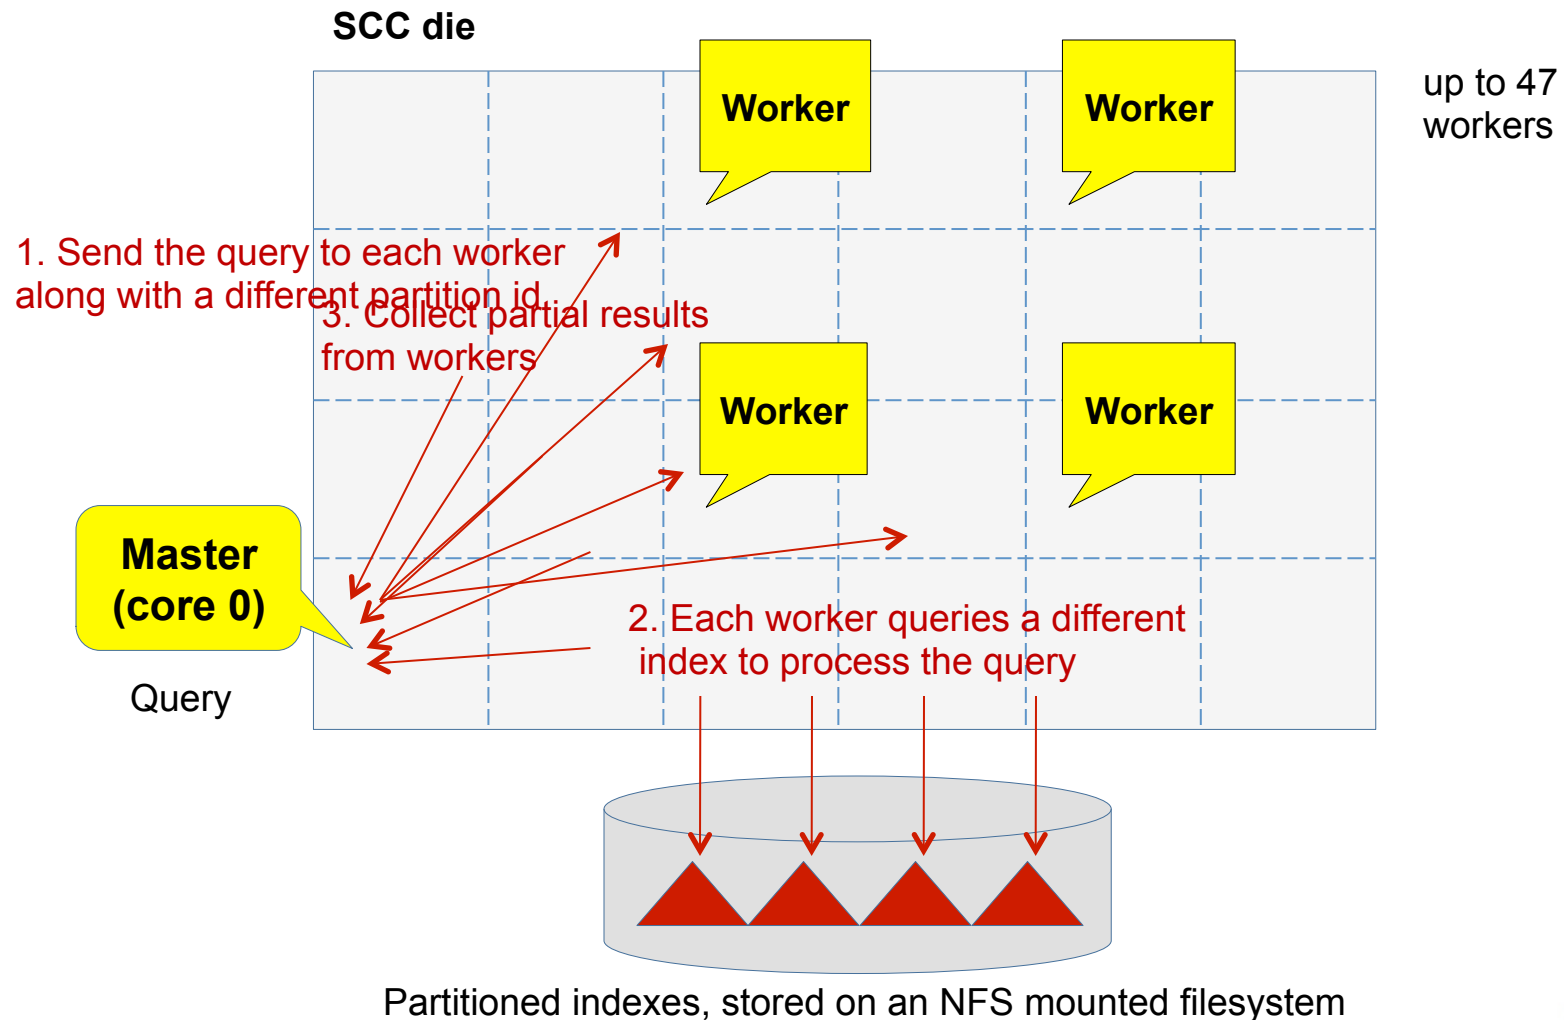
Models



Task Parallel Programming Model



Data Parallel Programming Model (1/2)



Data Parallel Programming Model (2/2)

- ◆ Partition the RDF graph into smaller graphs
 - ◆ Extract weakly connected directed subgraphs
 - ◆ Apply standard graph partitioning techniques (e.g., METIS [[SIAM '98](#)])
 - ◆ We may miss results
- ◆ Collect partial results
 - ◆ Multiple MPI_Recv
 - ◆ Single MPI_Gatherv

Performance Evaluation

- ◆ RDF-3X [[VLDB '08](#), [VLDB Journal '10](#)]
- ◆ RCKMPI (a modified MPICH2 for Intel SCC)
- ◆ Tile_Mesh_DDR: 800MHz, 800MHz, 800MHz
- ◆ 2GB index size (limit of OS)
- ◆ Indexes – stored on an NFS mounted filesystem
- ◆ Cold cache

Datasets

◆ Real datasets

- ◆ YAGO2 (27,331,797 triples)
- ◆ Uniprot (46,972,851 triples)

◆ Synthetic dataset

- ◆ LUBM (35,612,176 triples) [[WWW '05](#)]

◆ Data partitioning

- ◆ LUBM (based on RDF files)
- ◆ Uniprot (based on protein fragments)
- ◆ YAGO2
 - ◆ Star-shaped graphs
 - ◆ METIS

YAGO: Queries

| Query | I/O footprint | Type | % CPU | Serial time |
|-----------------|---------------|-------|-------|-------------|
| QY ₁ | 14,756 KB | small | 29 | 4.73 secs |
| QY ₂ | 15,004 KB | small | 40 | 9.23 secs |
| QY ₃ | 22,832 KB | small | 29 | 6.51 secs |
| QY ₄ | 33,492 KB | small | 21 | 9.27 secs |
| QY ₅ | 216,564 KB | large | 22 | 82.65 secs |
| QY ₆ | 272,848 KB | large | 30 | 120.08 secs |
| QY ₇ | 332,944 KB | large | 43 | 218.43 secs |

LUBM: Queries


| Query | I/O footprint | Type | % CPU | Serial time |
|-----------------|---------------|-------|-------|-------------|
| QL ₁ | 2,668 KB | small | 25 | 1.4 secs |
| QL ₂ | 3,132 KB | small | 35 | 1.47 secs |
| QL ₃ | 9,804 KB | small | 19 | 3.5 secs |
| QL ₄ | 636,204 KB | large | 32 | 299.99 secs |
| QL ₅ | 673,924 KB | large | 29 | 206.58 secs |

Uniprot: Queries

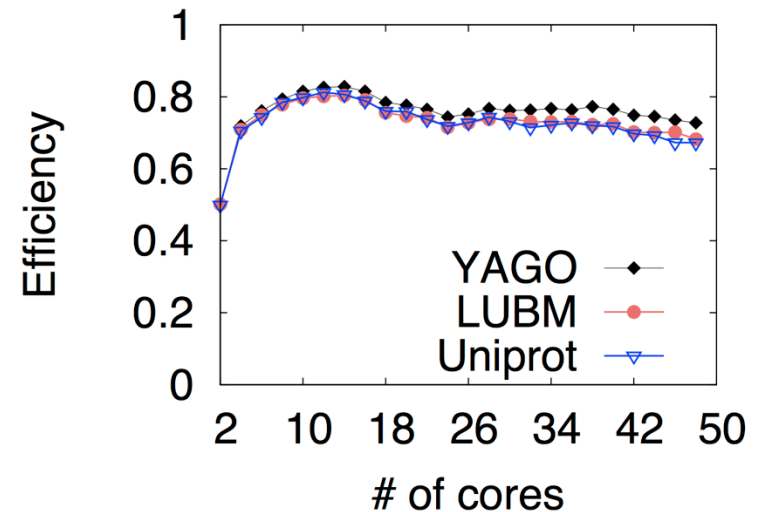
| Query | I/O footprint | Type | % CPU | Serial time |
|-----------------|---------------|-------|-------|-------------|
| QU ₁ | 4,468 KB | small | 39 | 2.08 secs |
| QU ₂ | 10,344 KB | small | 39 | 6.46 secs |
| QU ₃ | 48,020 KB | large | 31 | 19.39 secs |
| QU ₄ | 62,188 KB | large | 19 | 15.48 secs |
| QU ₅ | 166,808 KB | large | 17 | 43.51 secs |

ST Model (small I/O footprint, task parallelism)

Speedup

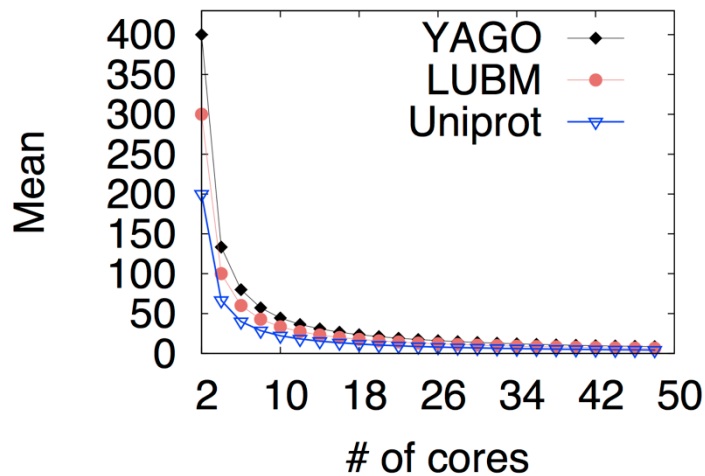


Efficiency

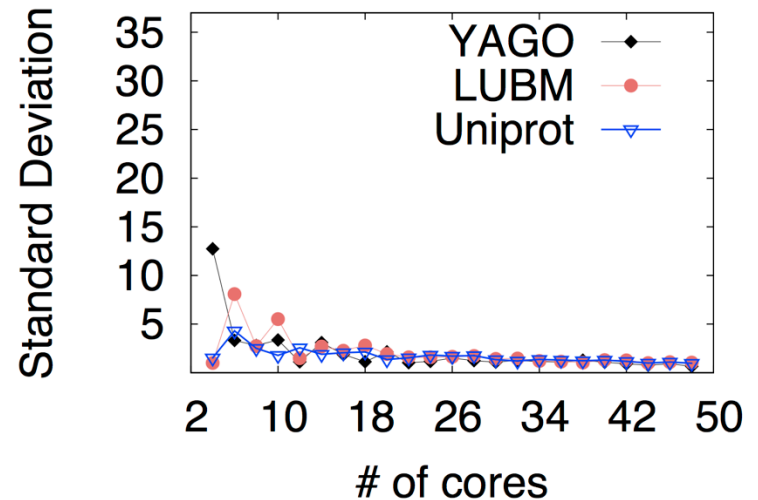


ST Model Load Distribution

Mean (# of tasks/worker)

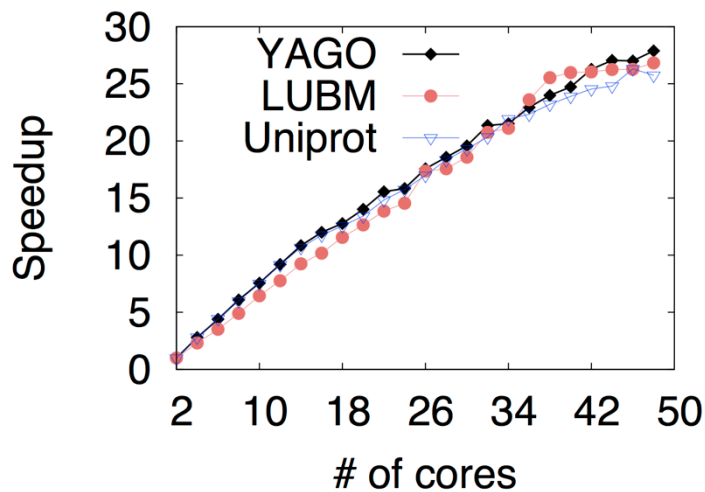


Standard deviation

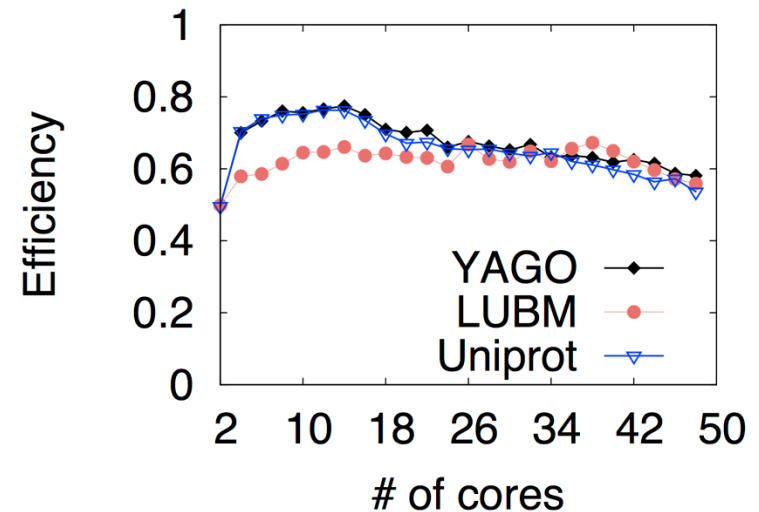


LT Model (large I/O footprint, task parallelism)

Speedup

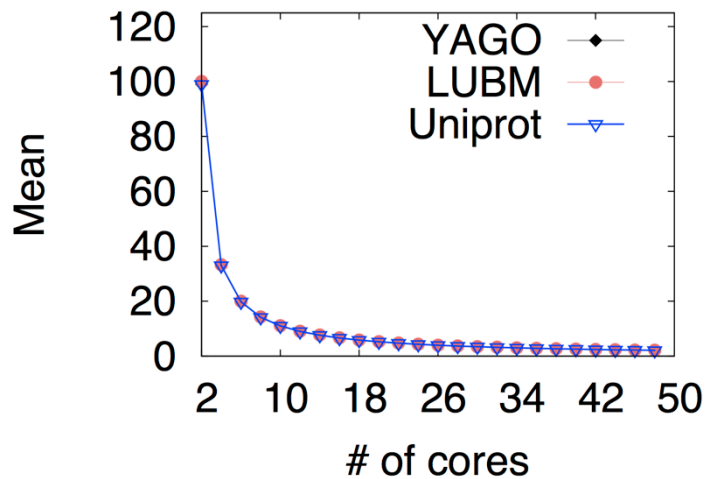


Efficiency

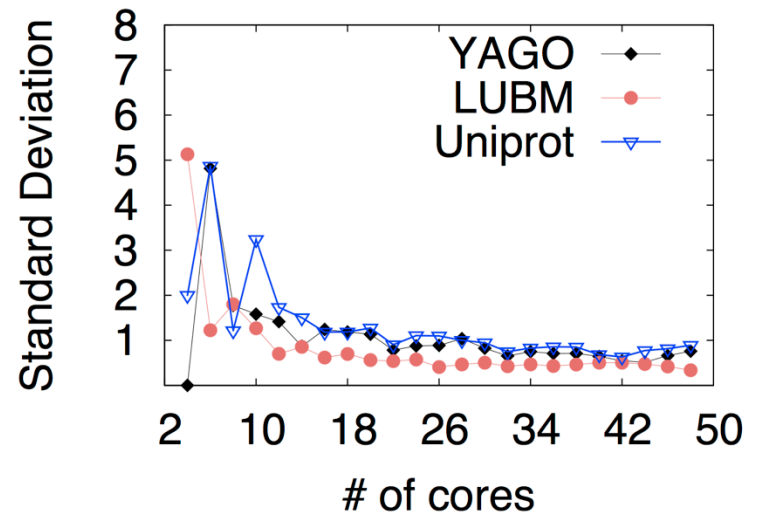


LT Model Load Distribution

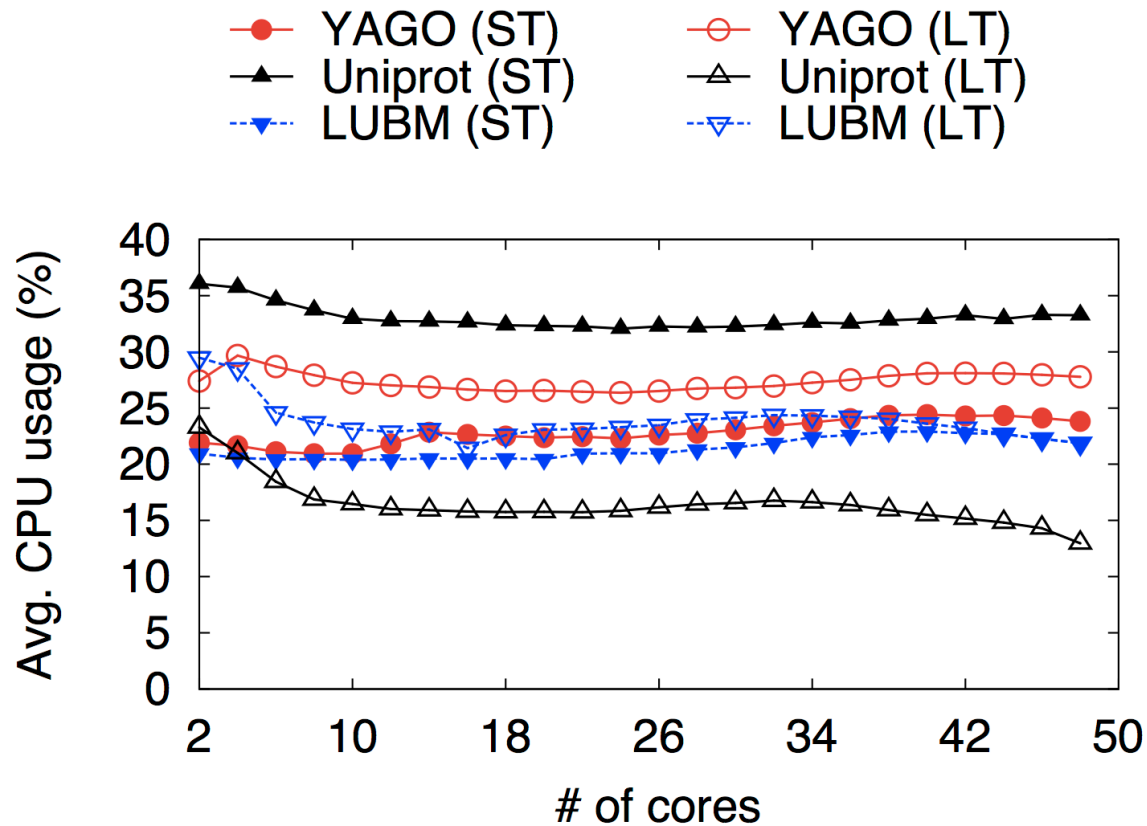
Mean (# of tasks/worker)



Standard deviation

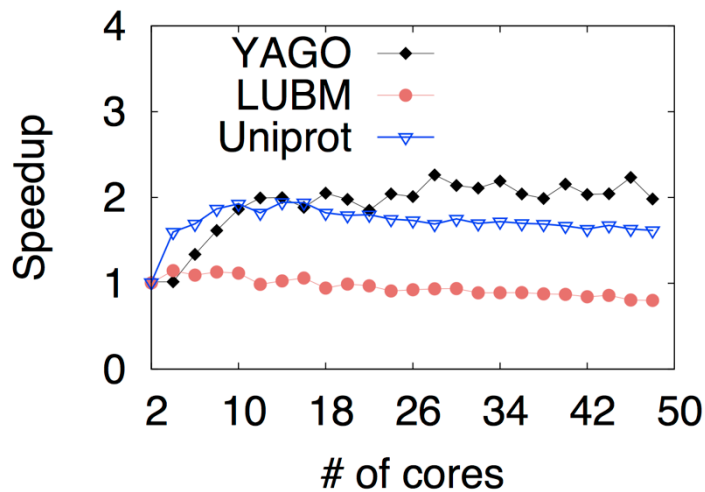


CPU Usage

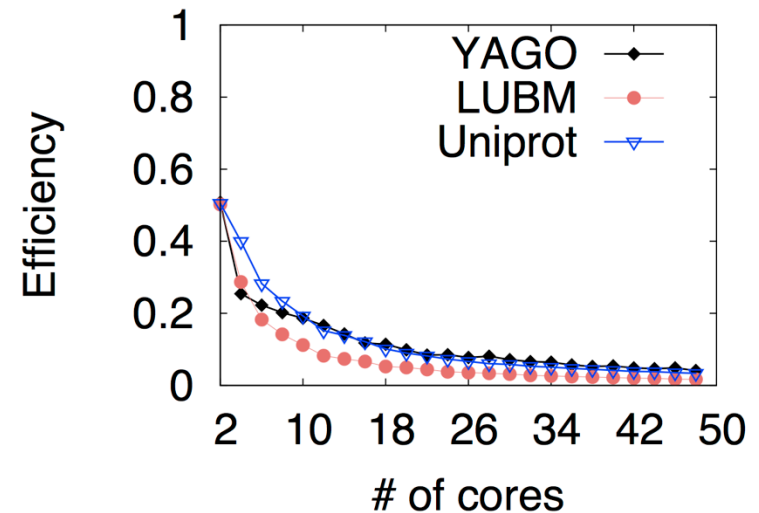


SD Model (small I/O footprint, data parallelism)

Speedup




Efficiency




LD Model (large I/O footprint, data parallelism)

Speedup



Efficiency



Conclusion and Future Work

- ◆ Task parallel programming yields good speedup and efficiency
 - ◆ For both large I/O and small I/O footprint queries
- ◆ Data parallel programming yields poor speedup and efficiency due to
 - ◆ Load imbalance or I/O contention
- ◆ Future work
 - ◆ New methods to overcome the challenges posed by the data parallel programming model
 - ◆ Effect of dynamic voltage and frequency scaling on the performance of RDF query processing

Questions?

◆ Acknowledgements

- ◆ Single-chip Cloud Computer Research Program, Intel Labs
- ◆ National Science Foundation (IIS-1115871), 2011-2014
- ◆ For further information, contact Prof. Praveen Rao (raopr@umkc.edu)