

A STUDY OF GOSSIP ALGORITHMS FOR INTERNET-SCALE CARDINALITY ESTIMATION OF DISTRIBUTED XML DATA

Vasil G. Slavov

Computer Science & Electrical Engineering
University of Missouri-Kansas City

June 22nd 2012

Acknowledgements

National Science Foundation (IIS-1115871)

Thesis Committee

- ◆ Praveen R. Rao, Ph.D., Committee Chair
- ◆ Yugyung Lee, Ph.D.
- ◆ Deep Medhi, Ph.D.

Overview

- ◆ Introduction
- ◆ Background
- ◆ Thesis objectives
- ◆ Design
 - ◆ VanillaXGossip
 - ◆ XGossip
- ◆ Implementation
- ◆ Performance evaluation
 - ◆ Amazon Elastic Compute Cloud (EC2)
- ◆ Conclusions

Introduction

◆ XML and XPath – W3C standards

```

<ClinicalDocument>
  <typeId extension="POCD_HD000040" root="2.16.840.1.113883.1.3"/>
  <id extension="CSE001" root="2.16.840.1.113883.19.4"/>
  <code code="8647-0" codeSystem="2.16.840.1.113883.6.1" codeSystemName="LOINC" displayName="Hospital Consultations"/>
  <confidentialityCode code="V" codeSystem="2.16.840.1.113883.5.25"/>
  - <RecordTarget>
    - <PatientRole>
      <ID>711</ID>
      <Patient>F</Patient>
      <ProviderOrganization>id root=2.16.840.1.113883.19.5</ProviderOrganization>
    </PatientRole>
  </RecordTarget>
  - <Author>
    <representedOrganization>UMKC School of Computing & Engineering (Research)</representedOrganization>
  </ClinicalDocument [RecordTarget/PatientRole/Patient = "M"] [RecordTarget/PatientRole/ID]
  - <section>
    <code code="" codeSystem="2.16.840.1.113883.6.1" codeSystemName="LOINC"/>
    <title>History of Present Illness</title>
    - <text>
      The patient was a 108-year-old nursing home resident , who was admitted with a two-day history of increased respiratory
      secretions and a 24-hour history of elevated fever . Despite Augmentin , the patient 's delirium worsened in the 24 hours prior to
      admission , and her temperature was up to 102 . She was refusing to take p.o.'s
    </text>
  </section>
  - <section>
    <code code="11348-0" codeSystem="2.16.840.1.113883.6.1" codeSystemName="LOINC"/>
    <title>Past Medical History</title>
  </section>

```

XPath

The Story So Far...

Galanis et.al. [VLDB '03],
 XPeer [P2P&DB '04],
 XP2P [WIDM '04],
 Garces et al. [ICDCS '04],
 Skobeltsyn et.al. [ODBASE'05],
 KadoP [ICDE '08],
 XTreeNet [VLDB '08],
 psiX [TKDE '09, ICDE '09],
 ...

Peer-to-peer
 Computing
 (Distributed Hash
 Tables)

Chord [SIGCOMM '01],
 CAN [SIGCOMM '01],
 Pastry [Middleware '01],
 Tapestry [JSAC '04],
 Kademlia [IPTPS '02]
 Dynamo [SOSP '07],
 Cassandra [SIGMOD '08],
 Voldemort [ICDE '11],
 ...

XML \cap P2P \cap Gossip

Gossip (or
 epidemic)
 algorithms

et. al. [PODC '87],
 Karp et. al. [FOCS '00],
 Kempe et.al. [FOCS '03],
 Berger et.al. [SODA '05],
 Ganesh et. al. [INFOCOMM '05],
 Boyd et. al. [INFOCOMM '05],
 Jelasity et. al. [TOCS '05],
 Kashyap et. al. [PODS '06],
 Georgiou et. al. [PODC '08],
 Mosk-Aoyama et. al. [TOIT '08],
 ...

XML \cap P2P = ?

Large-scale sharing of biomedical
and clinical data

```
graph TD; A[Large-scale sharing of biomedical and clinical data] --> B[Scalable clinical data sharing systems via a P2P architecture [Stead and Lin, 2009]]; A --> C[HL7 version 3 standard; XML based; semantic interoperability; can model discharge summaries, lab reports, ...];
```

Scalable clinical data sharing
systems via a P2P architecture
[[Stead and Lin, 2009](#)]

HL7 version 3 standard; XML based;
semantic interoperability; can model
discharge summaries, lab reports, ...

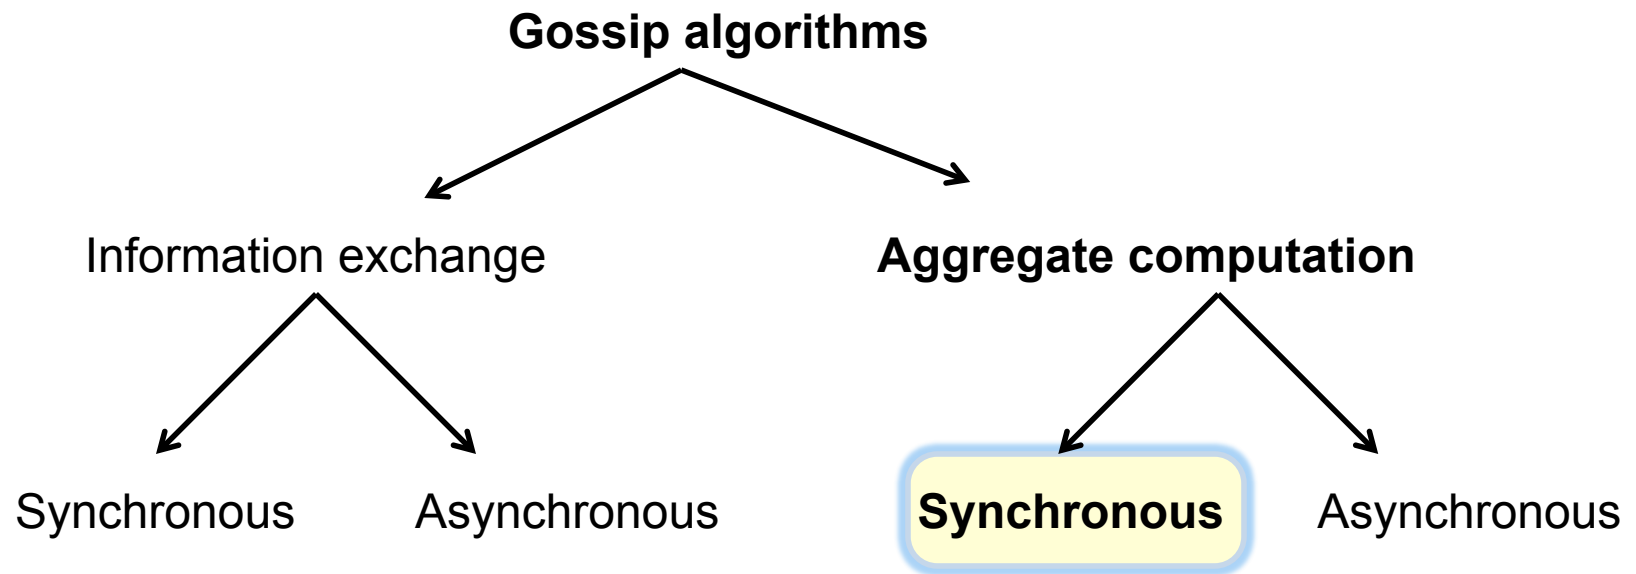
The Cancer Biomedical Informatics Grid (caBIG): a real world data sharing platform

The Problem: XPath Cardinality Estimation

- ◆ Compute the number of documents in the network that contain a match for the expression */Gene//goAcc*
- ◆ Useful for query optimization
- ◆ Desired properties
 - ◆ Scalability
 - ◆ Decentralization
 - ◆ Fault-tolerance
 - ◆ Efficient usage of bandwidth
 - ◆ Provable guarantee on the quality of the estimate

Gossip Algorithms

- ◆ Communication, computation, and information spreading
- ◆ Attractive in large-scale, distributed systems
- ◆ Real world examples: Amazon S3, Dynamo, Cassandra



Can further classify based on the topology of the network

Push-Sum Protocol


- ◆ By Kempe, Dobra, and Gehrke [FOCS '03]
- ◆ Each peer wishes to know the average
- ◆ Each peer maintains a (sum, weight) pair during gossip
 - ◆ In round @ $t = 0$: send $(f_1, 1)$ to itself
 - ◆ In any round @ $t > 0$: add up sums & weights, send $(s_1, s_2)/2, (w_1, w_2)/2$ to itself and to random peer
- ◆ Convergence (n peers)
 - ◆ Rounds: $O(\log(n) + \log(1/\epsilon) + \log(1/\delta))$
 - ◆ Messages: n messages per round
- ◆ Proof is based on the property of “mass conservation”

Thesis Objectives

1. Implementing gossip in an Internet-scale environment
2. Conducting a comprehensive evaluation
3. Analyzing the experimental results

Design of VanillaXGossip


- ◆ Builds on Push-Sum
- ◆ XML documents are mapped to their signatures
 - ◆ *psiX* [Rao et.al. TKDE '09, ICDE '09]
 - ◆ XML doc. \rightarrow data signature; XPath query \rightarrow query signature



Each peer creates a sorted tuple list

Merging Phase

- ◆ Suppose a peer receives 3 tuple lists during a gossip round



VanillaXGossip

◆ Special multiset \perp

- ◆ Placeholder for signatures not yet known to a peer during a gossip round
- ◆ Preserves the property of “mass conservation”

◆ Convergence

- ◆ Rounds: $O(\log(n) + \log(1/\epsilon) + \log(1/\delta))$

Problem ☹

A peer will end up with all the distinct signatures

More memory, more bandwidth

XGossip

♦ Idea

- ♦ **Divide-and-conquer** approach using Locality Sensitive Hashing (LSH)
- ♦ A **subset of peers** are responsible for gossiping a **subset of distinct signatures**

Benefits ☺


Less memory, less bandwidth, faster convergence

Locality Sensitive Hashing (LSH)

- ◆ Introduced by Indyk and Motwani [[STOC '98](#)]
- ◆ Applications
 - ◆ Web clustering, computer vision, computational biology, etc.
- ◆ Idea
 - ◆ Use many hash functions
 - ◆ Probability of collision is higher for inputs that are more similar
- ◆ LSH on sets using Jaccard index [[WWW '02](#), [WWW '05](#)]
 - ◆ $P[h(s_1) = h(s_2)] = |s_1 \cap s_2| / |s_1 \cup s_2|$
 - ◆ $h() \rightarrow$ min-hashing

LSH on Sets

- ◆ Suppose $\mathbf{p} = |s_1 \cap s_2| / |s_1 \cup s_2|$
- ◆ Pick $k \times \ell$ random linear hash functions



$P[\text{at least one } p]$

Can pick k and ℓ s.t.
 High probability if similarity $\geq \mathbf{p}$
 Low probability if similarity $< \mathbf{p}$

$$= 1 - (1 - \mathbf{p}^\ell)^k$$


XGossip (1/2)

- ◆ Define k teams for a signature \mathbf{s}

- ◆ $\text{LSH}(\mathbf{s}) \rightarrow \{h_1, \dots, h_k\}$

- ◆ Each team has id h_i , $1 \leq i \leq k$; Δ denotes team size

- ◆ $\alpha = 1 - (1 - \mathbf{p}^l)^k$



● denotes a peer

Cardinality estimation: more likely to find all the required signatures in the same team

XGossip (2/2)

- ◆ Initialization and execution phases
- ◆ Convergence
 - ◆ Rounds: $O(\log(\Delta) + \log(1/\varepsilon) + \log(1/\delta))$
- ◆ Bandwidth reduction
 - ◆ Compress signatures when sending a gossip message

VanillaXGossip: Cardinality Estimation

- ◆ Suppose a peer wants to compute $\text{card}(\text{/Gene//goAcc})$

Tuple list T @ the peer

Signature	(sum,weight)
s_i	(f_i, w_i)

Sum the frequency estimates of signatures that are supersets of the query signature; multiply by n

XGossip: Cardinality Estimation


- ◆ Suppose a peer wants to compute $\text{card}(\text{/Gene//goAcc})$

Apply LSH on the query signature $\rightarrow k$ teams

$T @ p_1$



$T @ p_2$



.....

$T @ p_k$

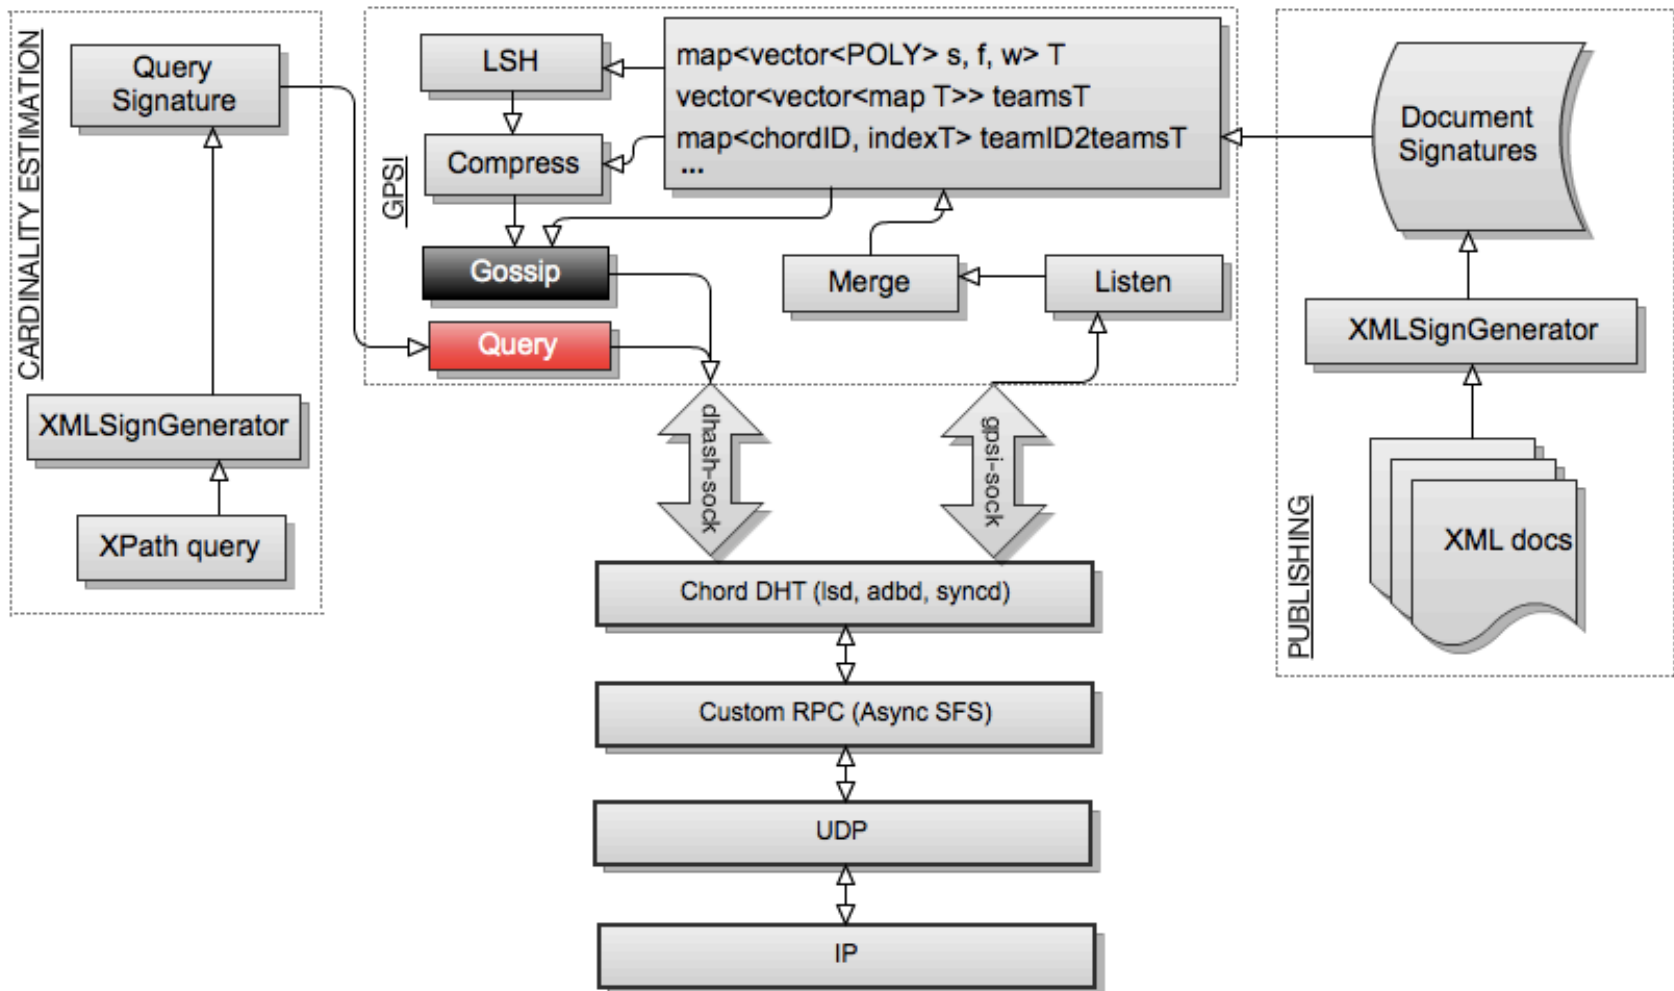


Merge the frequency estimates of signatures that are supersets of the query signature; multiply by Δ

Implementation

- ◆ Built on top of the Chord DHT [[SIGCOMM '01](#)]
 - ◆ Use Chord for routing (key-value pairs)
 - ◆ 4 processes per peer: Chord (lsd, syncd, adbd), Gossip (gpsi)
 - ◆ Communicate over UNIX sockets
 - ◆ Read signatures from files, store in main memory
- ◆ Implemented in C++
 - ◆ Data structures from STL
 - ◆ `typedef std::map<std::vector<POLY>, std::vector<double>, CompareSig> mapType;`
 - ◆ `typedef std::map<chordID, std::vector<int> > teamid2totalT;`
 - ◆ SFSlite asynchronous library
- ◆ Challenges

System Architecture



Performance Evaluation

◆ Datasets

- ◆ Generated by the IBM Synthetic XML generator using well-known DTDs (Treebank, DBLP, etc.)
- ◆ Uniformly distributed among all peers

Dataset	# of DTDs	Avg # of docs per DTD	Total # of docs	Avg. document signature size
D1	11	190,809	2,098,900	114 bytes
D2	13	192,223	2,498,900	127 bytes

Performance Evaluation

- ♦ Query sets
 - ♦ XPath queries generated by YFilter [[TODS '03](#)]
 - ♦ A total of 753 queries
 - ♦ 2 query approaches
 - ♦ LSH(XPath sig)
 - ♦ LSH(proxy sig)
 - ♦ Query subsets based on p_{\min} value

Query Set	Value of p_{\min}	# of queries
Q_0	$[0, 0.5)$	101
Q_1	$[0.5, 1]$	652
Q_2	$[0.6, 1]$	356
Q_3	$[0.7, 1]$	300
Q_4	$[0.8, 1]$	277
Q_5	$[0.9, 1]$	26

Network Setup and Distribution of Documents

- ◆ Amazon EC2
 - ◆ 20 medium instances (2 cores, 1.7GB memory)
 - ◆ US East availability zone
- ◆ Peers use local clock
- ◆ 120s long rounds
- ◆ Variables
 - ◆ team size (Δ): 8, 16
 - ◆ LSH parameter k : 4, 8
 - ◆ LSH parameter l : 10 (fixed): $\alpha = 0$, when $p < 0.5$
 - ◆ compression
 - ◆ # of peers: 500, 1000, 2000

Total # of peers in the network (n)	# of peers per EC2 instance	# of peers picked per DTD (z)
500	25	250
1000	50	500
2000	100	1000


Evaluation Metrics

- ◆ Convergence speed of the frequency of signatures
 - ◆ Mean absolute relative error (MARE) of the frequency estimate of the document signatures

$$\frac{1}{M} \times \sum_{i=1}^M \frac{|ef_i - f_i|}{f_i}$$


- ◆ Where
 - ◆ M: tuple list size
 - ◆ f_i : true signature frequency
 - ◆ ef_i : estimated signature frequency
 - ◆ VanillaXGossip: $\text{freq} / \text{weight} * n$
 - ◆ XGossip: $\text{freq} / \text{weight} * \Delta$
- ◆ Accuracy of cardinality estimation
 - ◆ MARE of the cardinality estimate of the queries
- ◆ Bandwidth consumption during gossip
 - ◆ Amount of data transmitted per round by all peers

XGossip in the Cloud




Diffusion Speed of Signatures

VanillaXGossip




Convergence of Frequencies of Signatures

VanillaXGossip



Dataset D_1 , $n = 1000$

XGossip

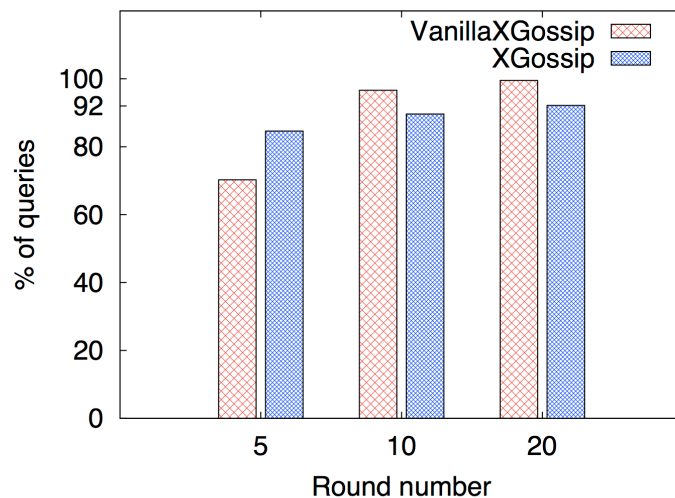


Dataset D_1 , $n = 1000$
 $\Delta = 8$, $k = 8$, $l = 10$

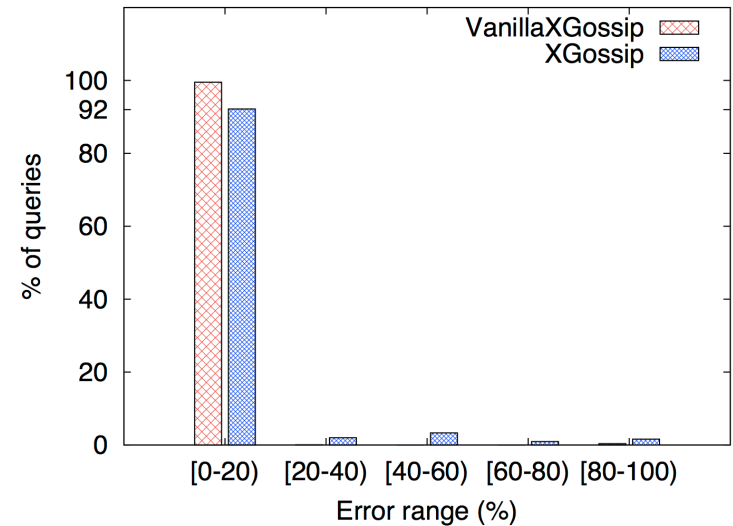
Accuracy of Cardinality Estimation (1/6)

VanillaXGossip vs. XGossip

At different rounds



After 20 rounds

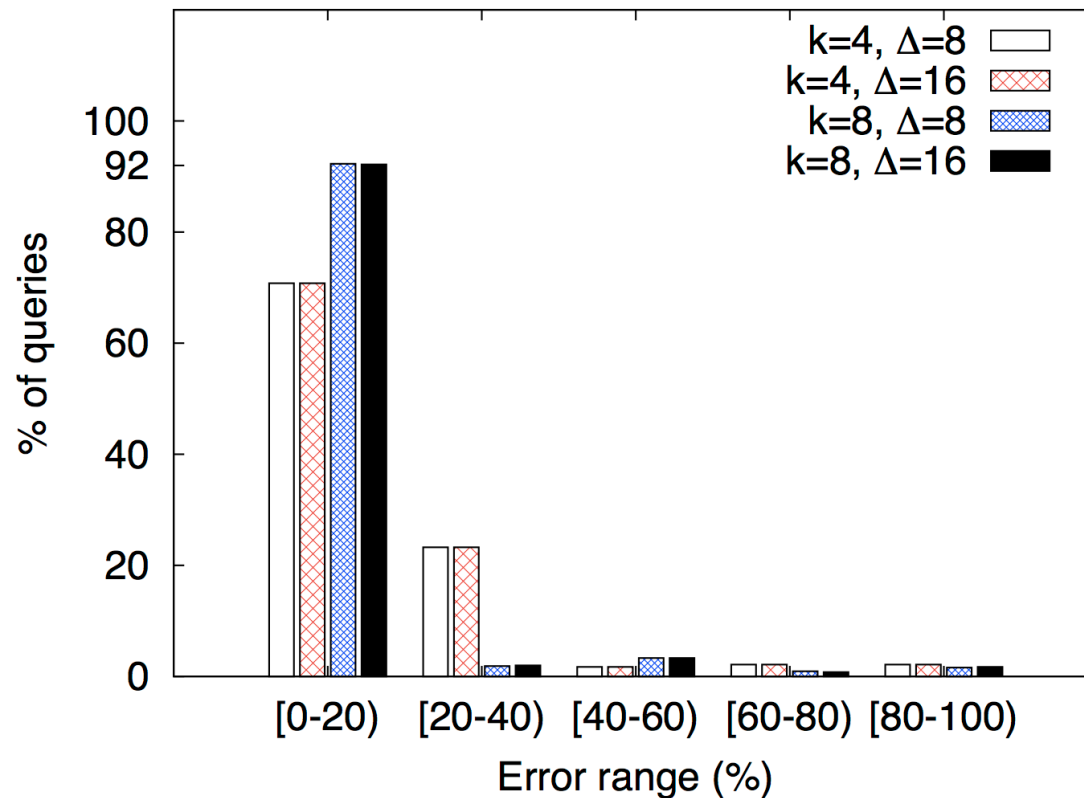


Relative error below 20%

$n = 1000$, $\Delta = 8$, $k = 8$, $l = 10$

Accuracy of Cardinality Estimation (2/6)

XGossip: LSH and team size




After 30 rounds, $n = 1000$, $l = 10$

Accuracy of Cardinality Estimation (3/6)


XGossip

$\Delta = 8, k = 4$



$n = 1000$, after 30 rounds
 $\Delta = 16, k = 4$ is almost identical

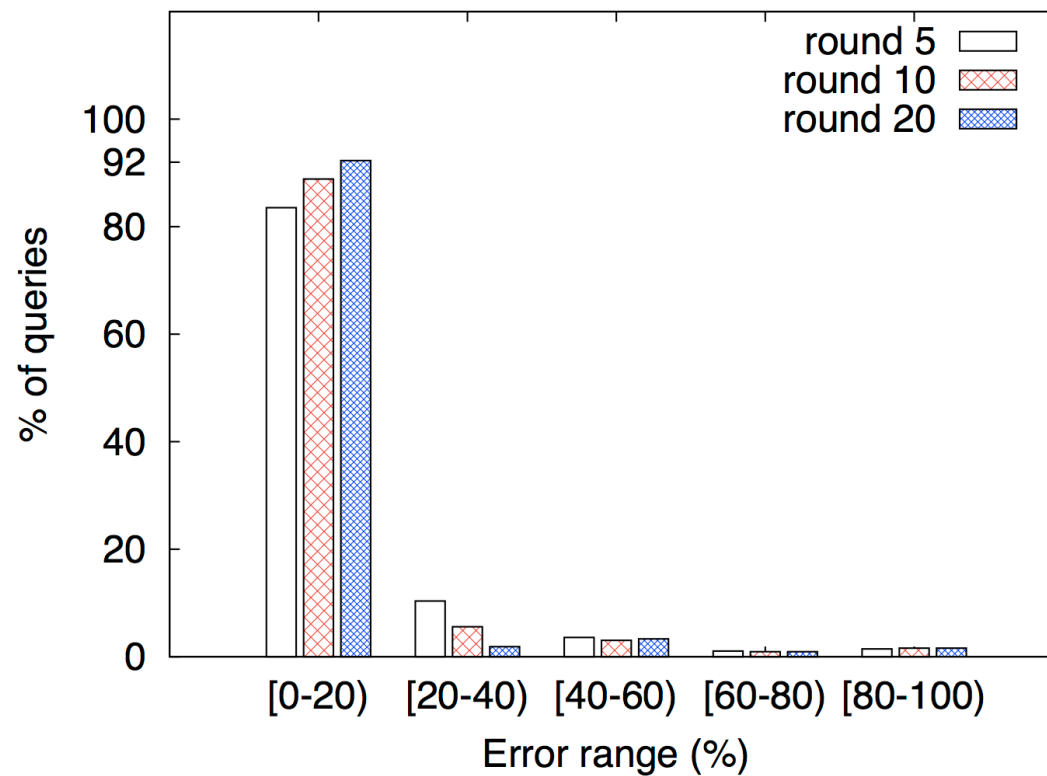
$\Delta = 8, k = 8$



$n = 1000$, after 30 rounds
 $\Delta = 16, k = 8$ is almost identical

Accuracy of Cardinality Estimation (4/6)

XGossip: increasing # of rounds

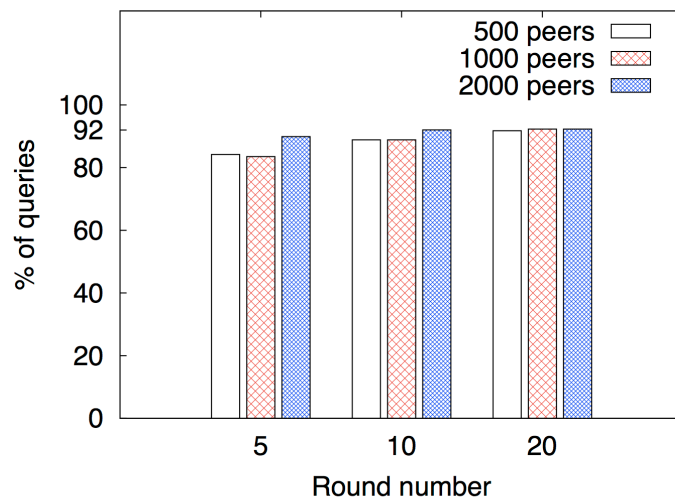


$n = 1000, \Delta = 8, k = 8, l = 10$

Accuracy of Cardinality Estimation (5/6)

XGossip


Varying # of peers



Relative error below 20%

$\Delta = 8, k = 8, l = 10$

Varying # of peers (query set Q_3 : [0.7, 1])



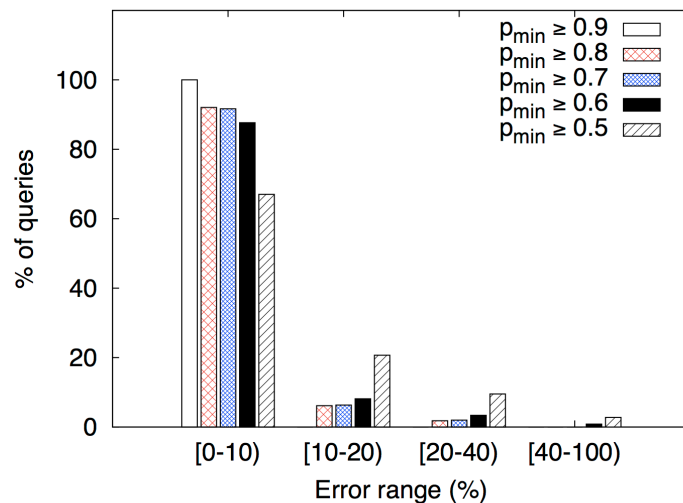
Relative error below 10%

$n = 1000, \Delta = 8, k = 8, l = 10$


Accuracy of Cardinality Estimation (6/6)

XGossip

500 peers



1000 peers




After 5 rounds

$\Delta = 8, k = 8, l = 10$


Bandwidth Consumption (1/2)

VanillaXGossip vs. XGossip (Dataset D_1)



VanillaXGossip: 10,309 MB
XGossip: 484 MB

XGossip compression (Dataset D_2)




XGossip (no compression): 9,874.2 MB
XGossip: 1,805.9 MB

Bandwidth Consumption (2/2)


XGossip

Different values of k and Δ



$n = 1000$

Varying the # of peers



$l = 10$

Performance Analysis: Results Summary

- ◆ LSH: tune k and l
- ◆ Compression works
- ◆ XGossip scales, VanillaXGossip does not
- ◆ XGossip converges faster than VanillaXGossip
- ◆ XGossip transmits less data than VanillaXGossip

# of peers	Avg. # of teams/peer	Avg. # of sigs/peer	Avg. # of sigs/team	Avg. msg size/peer (bytes)	Total # of msgs
500	88.40	4,024.25	45.52	1,160.18	440,240
1000	44.82	2,040.81	45.52	1,265.64	440,240
2000	23.09	1051.20	45.52	1,244.91	441,750

Conclusion

◆ Thesis objectives

1. Implementing gossip in an Internet-scale environment
2. Conducting a comprehensive evaluation
3. Analyzing the experimental results

◆ The results we obtained were consistent with the theoretical analysis of VanillaXGossip and XGossip.

Questions?

◆ References

- ◆ Vasil G. Slavov, Praveen R. Rao - **Towards Internet-Scale Cardinality Estimation of XPath Queries over Distributed XML Data**. *Proceedings of 6th International Workshop on Networking Meets Databases (NetDB 2011)*, Athens, Greece.
- ◆ Praveen Rao and Vasil Slavov - **Towards Internet-Scale Cardinality Estimation of XPath Queries over Distributed XML Data**. University of Missouri-Kansas City, Kansas City, MO 64110, Tech. Rep. TR-DB-2011-01, Jun. 2011, <http://r.faculty.umkc.edu/raopr/TR-DB-2011-01.pdf>.

◆ Acknowledgements

- ◆ National Science Foundation (IIS-1115871), 2011-2014