

INTERNET-SCALE CARDINALITY ESTIMATION OF XPATH QUERIES OVER DISTRIBUTED XML DATA

Praveen R. Rao and Vasil G. Slavov

Computer Science & Electrical Engineering

University of Missouri-Kansas City

Acknowledgements

National Science Foundation (IIS-1115871), University of Missouri Research Board, IBM Smarter Planet Faculty Innovation Award 2010

Roadmap

- ◆ Introduction
- ◆ Background
- ◆ XPath cardinality estimation
 - ◆ VanillaXGossip
 - ◆ XGossip
- ◆ Implementation and evaluation
 - ◆ Amazon Elastic Compute Cloud (EC2)
- ◆ Conclusions

By the end of this talk...

- ◆ “Gossip is good” in large-scale distributed systems



Introduction

◆ XML and XPath – W3C standards

```

<ClinicalDocument>
  <typeId extension="POCD_HD000040" root="2.16.840.1.113883.1.3"/>
  <id extension="CSE001" root="2.16.840.1.113883.19.4"/>
  <code code="8647-0" codeSystem="2.16.840.1.113883.6.1" codeSystemName="LOINC" displayName="Hospital Consultations"/>
  <confidentialityCode code="V" codeSystem="2.16.840.1.113883.5.25"/>
  - <RecordTarget>
    - <PatientRole>
      <ID>711</ID>
      <Patient>F</Patient>
      <ProviderOrganization>id root=2.16.840.1.113883.19.5</ProviderOrganization>
    </PatientRole>
  </RecordTarget>
  - <Author>
    <representedOrganization>UMKC School of Computing & Engineering (Research)</representedOrganization>
  </ClinicalDocument [RecordTarget/PatientRole/Patient = "M"] [RecordTarget/PatientRole/ID]
  - <section>
    <code code="" codeSystem="2.16.840.1.113883.6.1" codeSystemName="LOINC"/>
    <title>History of Present Illness</title>
    - <text>
      The patient was a 108-year-old nursing home resident , who was admitted with a two-day history of increased respiratory
      secretions and a 24-hour history of elevated fever . Despite Augmentin , the patient 's delirium worsened in the 24 hours prior to
      admission , and her temperature was up to 102 . She was refusing to take p.o.'s
    </text>
  </section>
  - <section>
    <code code="11348-0" codeSystem="2.16.840.1.113883.6.1" codeSystemName="LOINC"/>
    <title>Past Medical History</title>
  </section>

```

XPath

The Story So Far...

Galanis et.al. [VLDB '03],
 XPeer [P2P&DB '04],
 XP2P [WIDM '04],
 Garces et al. [ICDCS '04],
 Skobeltsyn et.al. [ODBASE'05],
 KadoP [ICDE '08],
 XTreeNet [VLDB '08],
 psiX [TKDE '09, ICDE '09],
 ...

Peer-to-peer
 Computing
 (Distributed Hash
 Tables)

Chord [SIGCOMM '01],
 CAN [SIGCOMM '01],
 Pastry [Middleware '01],
 Tapestry [JSAC '04],
 Kademlia [IPTPS '02]
 Dynamo [SOSP '07],
 Cassandra [SIGMOD '08],
 Voldemort [ICDE '11],
 ...

XML \cap P2P \cap Gossip

Gossip (or
 epidemic)
 algorithms

et. al. [PODC '87],
 Karp et. al. [FOCS '00],
 Kempe et.al. [FOCS '03],
 Berger et.al. [SODA '05],
 Ganesh et. al. [INFOCOMM '05],
 Boyd et. al. [INFOCOMM '05],
 Jelasity et. al. [TOCS '05],
 Kashyap et. al. [PODS '06],
 Georgiou et. al. [PODC '08],
 Mosk-Aoyama et. al. [TOIT '08],
 ...

XML \cap P2P = ?

Large-scale sharing of biomedical
and clinical data

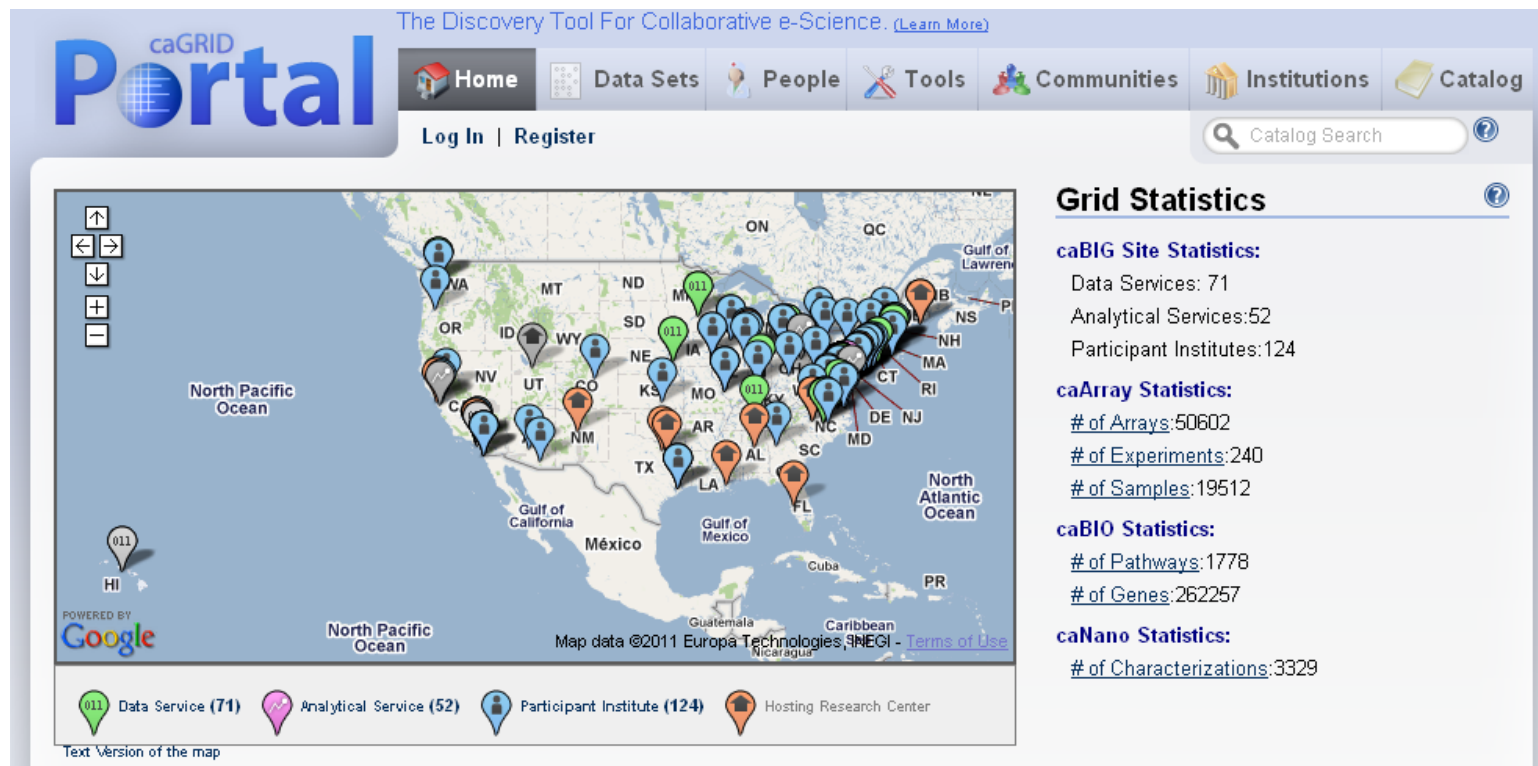
```
graph TD; A[Large-scale sharing of biomedical and clinical data] --> B[Scalable clinical data sharing systems via a P2P architecture [Stead and Lin, 2009]]; A --> C[HL7 version 3 standard; XML based; semantic interoperability; can model discharge summaries, lab reports, ...];
```

Scalable clinical data sharing
systems via a P2P architecture
[[Stead and Lin, 2009](#)]

HL7 version 3 standard; XML based;
semantic interoperability; can model
discharge summaries, lab reports, ...

A Real World Data Sharing Platform

- ◆ The Cancer Biomedical Informatics Grid (caBIG)
 - ◆ Source: <http://cagrid-portal.nci.nih.gov/>
 - ◆ About 124 participants across the US (SOA, XML databases)




Example from caBIG

“Find all the expression data where there are at least 50 conditions for genes found in the vacuole”

```
FOR $gene IN service
  ("http://cabio.osu.edu/GeneService.wsdl")/Gene,
$go IN service
  ("http://cabio.osu.edu/GeneOntologyService.wsdl")/GeneOntology,
$microarray IN service
  ("http://caarray.duke.edu/caArrayService.wsdl")/Microarray
LET $subject := $microarray/experiment/subject
WHERE
  $go/term='vacuole' AND $gene/goAcc=$go/acc AND
  $gene/gbAcc=$microarray/data/geneId AND
  count($microarray/data[geneId=$gene/$gbAcc]/condition)>50
RETURN
<subject>
  <subjectId>{ $subject/lsid }</subjectId>
  <species>{ $subject/species }</species>
  <microarrayData>
    { $microarray/data }
  </microarrayData>
</subject>
```


The Problem: XPath Cardinality Estimation



Compute the number of documents in the network that contain a match for the expression $/Gene/goAcc$

XPath Cardinality Estimation

◆ Useful for

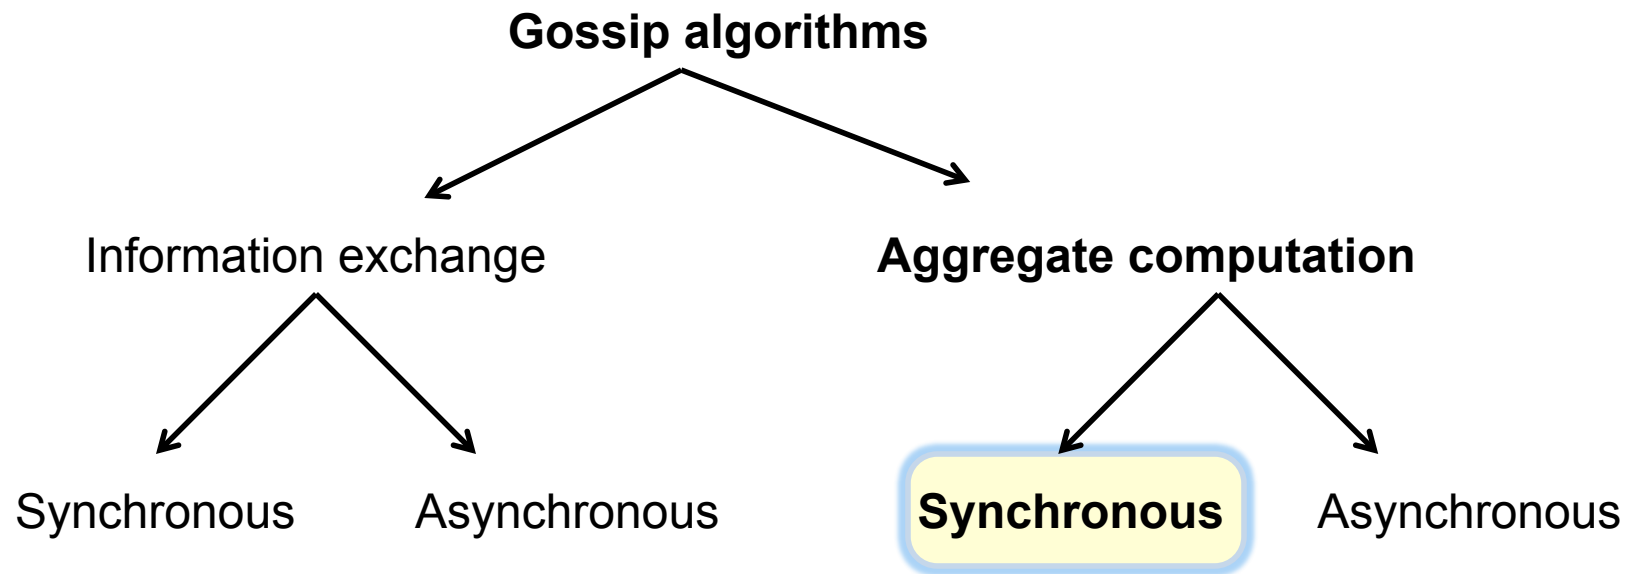
- ◆ XQuery optimization
 - ◆ E.g., to select a particular join ordering
- ◆ Designing IR-style ranking schemes
- ◆ Designing clinical studies
 - ◆ E.g., to determine if sufficient number of samples available to conduct a study

Desired properties

Scalability, decentralization, fault-tolerance, efficient usage of bandwidth, provable guarantee on the quality of the estimate

Gossip Algorithms

- ◆ Communication, computation, and information spreading
- ◆ Attractive in large-scale, distributed systems




Can further classify based on the topology of the network

Real World Examples

- ◆ Gossip algorithms are used in practice
 - ◆ Amazon's Dynamo (key-value store)
 - ◆ Amazon's S3 data centers
 - ◆ Facebook's Cassandra (key-value store)

Push-Sum Protocol (1/4)


- ◆ By Kempe, Dobra, and Gehrke [FOCS '03]
- ◆ Each peer wishes to know the average



$$\text{avg} = \frac{(f_1 + f_2 + f_3 + f_4)}{4}$$


Push-Sum Protocol (2/4)

- ◆ Each peer maintains a (sum, weight) pair during gossip
- ◆ In the round @ $t = 0$



Push-Sum Protocol (3/4)

♦ In any round @ $t > 0$




Proof is based on the property of “mass conservation”

Messages: n messages per round

VanillaXGossip


- ◆ Builds on Push-Sum
- ◆ XML documents are mapped to their signatures
 - ◆ *psiX* [Rao et.al. TKDE '09, ICDE '09]
 - ◆ XML doc. \rightarrow data signature; XPath query \rightarrow query signature



Each peer creates a sorted tuple list

Merging Phase

- ◆ Suppose a peer receives 3 tuple lists during a gossip round



VanillaXGossip

- ◆ Special multiset \perp
 - ◆ Placeholder for signatures not yet known to a peer during a gossip round
 - ◆ Preserves the property of “mass conservation”
- ◆ Convergence
 - ◆ Rounds: $O(\log(n) + \log(1/\epsilon) + \log(1/\delta))$

Problem ☹

A peer will end up with all the distinct signatures

More memory, more bandwidth

XGossip

♦ Idea

- ♦ **Divide-and-conquer** approach using Locality Sensitive Hashing (LSH)
- ♦ A **subset of peers** are responsible for gossiping a **subset of distinct signatures**

Benefits ☺


Less memory, less bandwidth, faster convergence

Locality Sensitive Hashing (LSH)

- ◆ Introduced by Indyk and Motwani [[STOC '98](#)]
- ◆ Applications
 - ◆ Web clustering, computer vision, computational biology, etc.
- ◆ Idea
 - ◆ Use many hash functions
 - ◆ Probability of collision is higher for inputs that are more similar
- ◆ LSH on sets using Jaccard index [[WWW '02](#), [WWW '05](#)]
 - ◆ $P[h(s_1) = h(s_2)] = |s_1 \cap s_2| / |s_1 \cup s_2|$
 - ◆ $h() \rightarrow$ min-hashing

LSH on Sets

- ◆ Suppose $\mathbf{p} = |s_1 \cap s_2| / |s_1 \cup s_2|$
- ◆ Pick $k \times \ell$ random linear hash functions




$P[\text{at least one } p]$

Can pick k and ℓ s.t.
 High probability if similarity $\geq \mathbf{p}$
 Low probability if similarity $< \mathbf{p}$

$$= 1 - (1 - \mathbf{p}^\ell)^k$$

XGossip (1/2)

- ◆ Define k teams for a signature \mathbf{s}
 - ◆ $\text{LSH}(\mathbf{s}) \rightarrow \{h_1, \dots, h_k\}$
 - ◆ Each team has id h_i , $1 \leq i \leq k$; Δ denotes team size



● denotes a peer

Cardinality estimation: more likely to find all the required signatures in the same team

XGossip (2/2)

- ◆ Initialization and execution phases
- ◆ Convergence
 - ◆ Rounds: $O(\log(\Delta) + \log(1/\varepsilon) + \log(1/\delta'))$
- ◆ Bandwidth reduction
 - ◆ Compress signatures when sending a gossip message

VanillaXGossip: Cardinality Estimation

- ◆ Suppose a peer wants to compute $\text{card}(\text{/Gene//goAcc})$

Tuple list T @ the peer

Signature	(sum,weight)
s_i	(f_i, w_i)

Sum the frequency estimates of signatures that are supersets of the query signature; multiply by n

XGossip: Cardinality Estimation


- ◆ Suppose a peer wants to compute $\text{card}(\text{/Gene//goAcc})$

Apply LSH on the query signature $\rightarrow k$ teams

$T @ p_1$



$T @ p_2$



.....

$T @ p_k$



Merge the frequency estimates of signatures that are supersets of the query signature; multiply by Δ

Asymptotic Analysis

Metric	VanillaXGossip	XGossip
Accuracy	$r\epsilon$	$r\epsilon$
Confidence	$(1 - \delta)$	$(1 - \delta)$
Convergence (# of rounds)	$O(\log(n) + \log(\frac{1}{\epsilon}) + \log(\frac{1}{\delta}))$	$O(\log(\Delta) + \log(\frac{1}{\epsilon}) + \log(\frac{\alpha}{\alpha + \delta - 1}))$
Bandwidth	$O(nD)$	$O(\log(n)kD\Delta)$
Messages	$O(n \log(n))$	$O(\frac{\log(n)}{n}kD\Delta \log(\Delta))$

$r \rightarrow$ # of signatures that are supersets of a query signature

α depends on the minimum similarity between the query signature and the r distinct document signatures that it divides, and k and ℓ

$D \rightarrow$ # of distinct signatures in the network

A peer becomes a successor of $O(kD \log(n)/n)$ teams (using the property of consistent hashing)

of distinct signatures per team $\rightarrow O(kD\Delta \log(n)/n)$

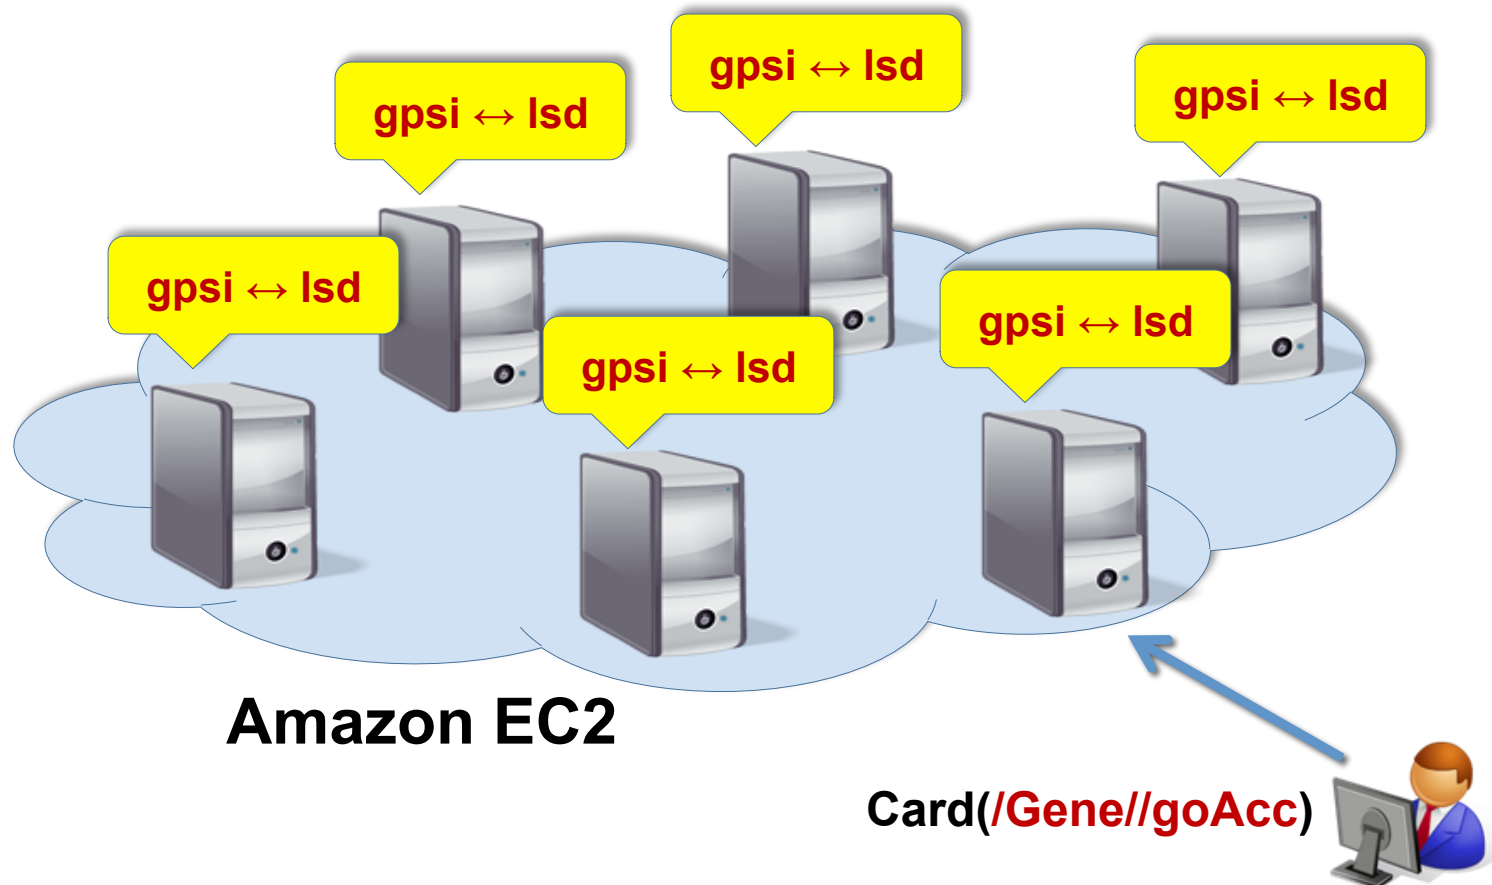
VanillaXGossip: will always find all the superset signatures

XGossip: may miss a superset signature if its similarity with the query signature is below the threshold used by LSH

Implementation

- ◆ Built on top of the Chord DHT [[SIGCOMM '01](#)]
 - ◆ Use Chord for routing (key-value pairs)
 - ◆ 4 processes per peer: Chord (lsd, syncd, adbd), Gossip (gpsi)
 - ◆ Communicate over UNIX sockets
 - ◆ Read signatures from files, store in main memory

System Architecture




Performance Evaluation

- ♦ Amazon EC2
 - ♦ 20 medium instances (2 cores, 1.7GB memory), 50 peers/instance, 1000 peers
- ♦ Datasets
 - ♦ Generated by the IBM Synthetic XML generator using well-known DTDs (Treebank, DBLP, etc.)
 - ♦ Uniformly distributed among all peers (25 hosts/DTD)
- ♦ Queries
 - ♦ XPath queries generated by YFilter [[TODS '03](#)]
 - ♦ LSH(XPath): VanillaXGossip, XGossip
 - ♦ LSH(proxy sig): XGossip
- ♦ Variables: team size (8, 16); LSH: K (4, 8), L (10); compression

Number of DTDs	Avg # of docs per DTD	Total # of docs	Avg. document signature size	Total # of queries
11	9,540	104,945	114 bytes	1977
13	9,611	124,945	127 bytes	2355

Query Error (1/3)


VanillaXGossip vs XGossip



- ◆ Frequency of queries with relative error within the specified ranges

Query Error (2/3)


VanillaXGossip vs XGossip: rounds



- ◆ Frequency of queries with relative error below 20%

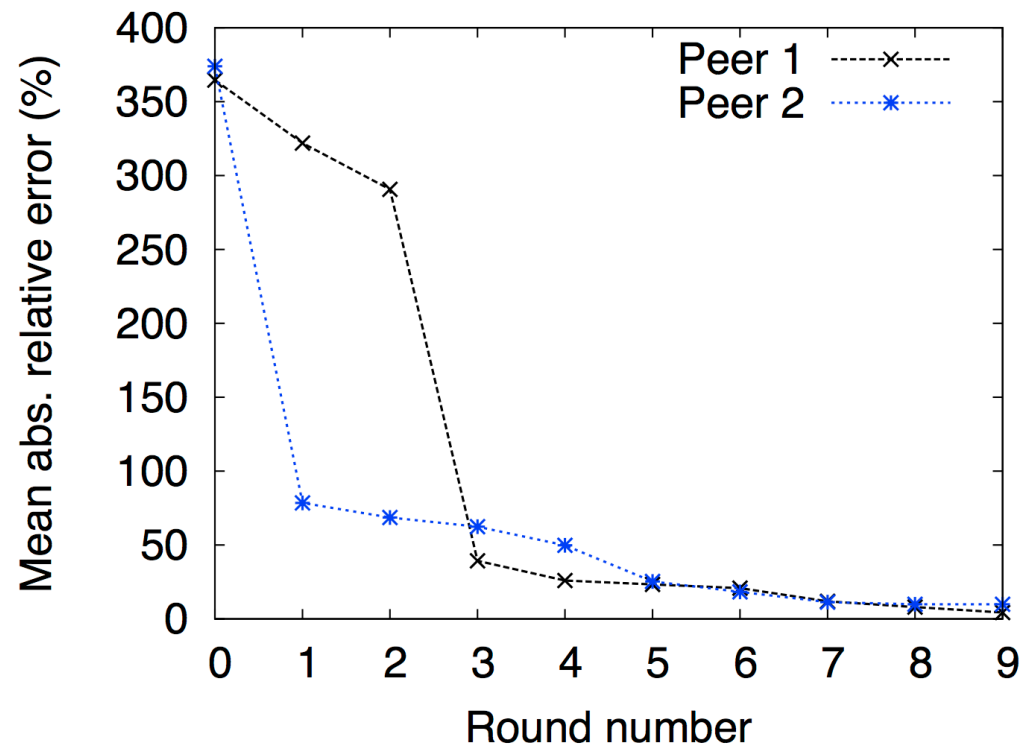
Query Error (3/3)

XGossip: LSH and team size



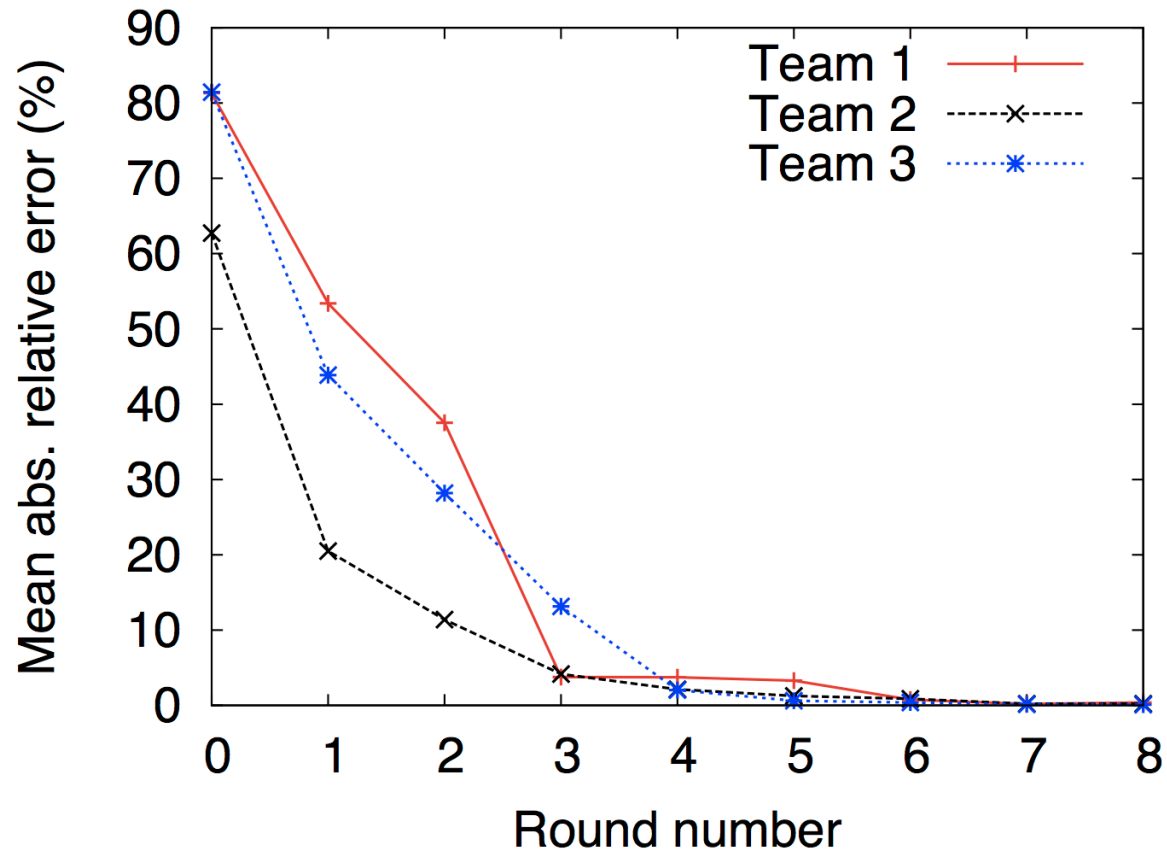
Signature Convergence (1/2)

VanillaXGossip



Signature Convergence (2/2)

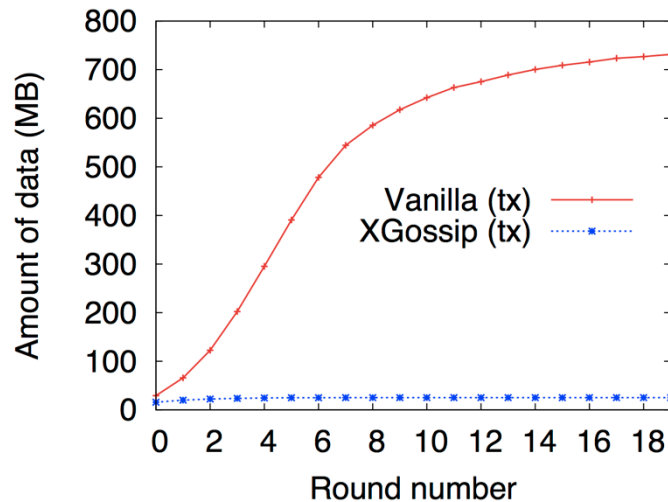
XGossip




Bandwidth Consumption (1/2)

VanillaXGossip vs. XGossip

Transmitted data



Received data

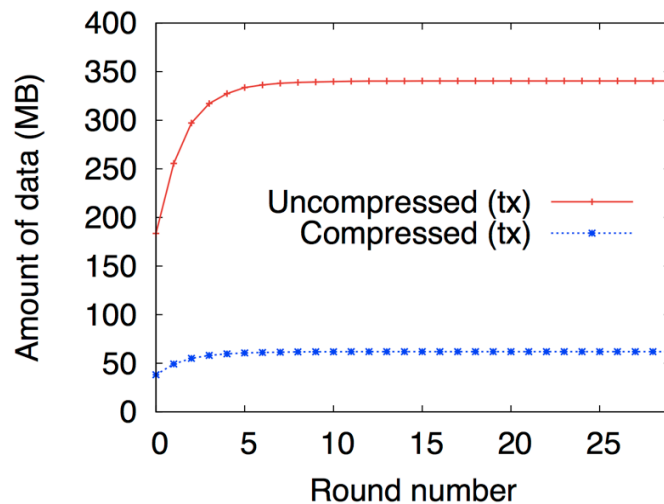


- ♦ VanillaXGossip: 10,309 MB
- ♦ XGossip: 484 MB

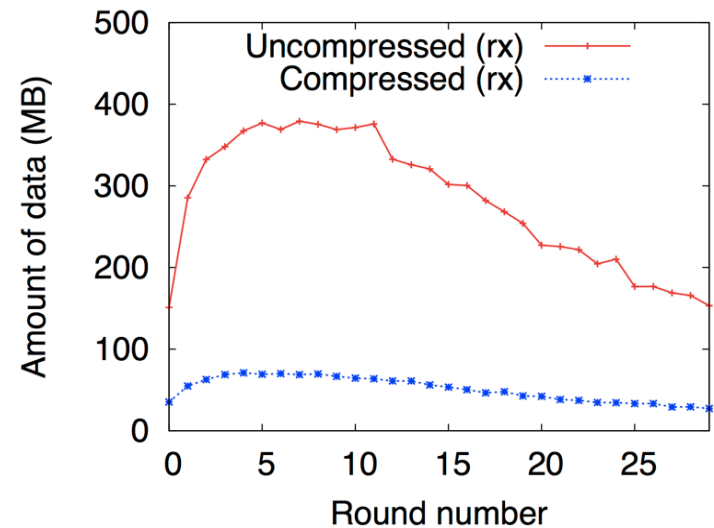
Bandwidth Consumption (2/2)

XGossip compression

Transmitted data



Received data



- ◆ Uncompressed: 9874 MB
- ◆ Compressed: 1805 MB

Questions?

◆ References

- ◆ Vasil G. Slavov, Praveen R. Rao - **Towards Internet-Scale Cardinality Estimation of XPath Queries over Distributed XML Data.** *Proceedings of 6th International Workshop on Networking Meets Databases (NetDB 2011)*, Athens, Greece.

◆ Acknowledgements

- ◆ National Science Foundation (IIS-1115871), 2011-2014
- ◆ University of Missouri Research Board, 2010-2012
- ◆ IBM Smarter Planet Faculty Innovation Award, 2010